

Contents

I	Preliminaries	
1	Foundations of Machine Learning	17
1.1	Math Basics	18
1.1.1	Linear Algebra	18
1.1.2	Probability and Statistics	22
1.2	Designing a Text Classifier	25
1.2.1	Problem Formulation	26
1.2.2	Documents as Feature Vectors	28
1.2.3	Linear Classifiers	29
1.2.4	Generative vs Discriminative	30
1.2.5	OOV Words and Smoothing	34
1.3	General Problems	36
1.3.1	Supervised and Unsupervised Models	36
1.3.2	Inductive Bias	37
1.3.3	Non-linearity	39
1.3.4	Training and Loss Functions	41
1.3.5	Overfitting and Underfitting	46
1.3.6	Prediction and Inference	48
1.4	Model Selection and Evaluation	49
1.4.1	Strategies for Model Selection	50
1.4.2	Training, Validation and Test Data	55
1.4.3	Performance Measure	56
1.4.4	Significance Tests	56
1.5	NLP Tasks as ML Tasks	58
1.5.1	Classification	58
1.5.2	Sequence Labeling	59
1.5.3	Language Modeling/Word Prediction	60
1.5.4	Sequence Generation	61
1.5.5	Tree Generation	62
1.5.6	Relevance Modeling	64

1.5.7	Linguistic Alignment	65
1.5.8	Information Extraction	66
1.5.9	Remarks	67
1.6	Summary	67
2	Foundations of Neural Networks	69
2.1	Multi-layer Neural Networks	69
2.1.1	Single-layer Perceptrons	69
2.1.2	Stacking Multiple Layers	71
2.1.3	Computation Graphs	73
2.2	Example: Neural Language Modeling	76
2.3	Basic Model Architectures	81
2.3.1	Recurrent Units	81
2.3.2	Convolutional Units	83
2.3.3	Gate Units	86
2.3.4	Normalization (Standardization) Units	86
2.3.5	Residual Units	87
2.4	Training Neural Networks	88
2.4.1	Gradient Descent	89
2.4.2	Batching	93
2.4.3	Parameter Initialization	94
2.4.4	Learning Rate Scheduling	95
2.5	Regularization Methods	98
2.5.1	Norm-based Penalties	98
2.5.2	Dropout	99
2.5.3	Early Stopping	101
2.5.4	Smoothing Output Probabilities	102
2.5.5	Training with Noise	104
2.6	Unsupervised Methods and Auto-encoders	106
2.6.1	Auto-encoders with Explicit Regularizers	109
2.6.2	Denoising Auto-encoders	112
2.6.3	Variational Auto-encoders	114
2.7	Summary	118

II

Basic Models

3	Words and Word Vectors	121
3.1	Tokenization	121
3.1.1	Tokenization via Rules and Heuristics	122
3.1.2	Tokenization as Language Modeling	124

3.1.3	Tokenization as Sequence Labeling	127
3.1.4	Learning Subwords	128
3.2	Vector Representation for Words	135
3.2.1	One-hot Representation	135
3.2.2	Distributed Representation	136
3.2.3	Compositionality and Contextuality	138
3.3	Count-based Models	140
3.3.1	Co-occurrence Matrices	140
3.3.2	TF-IDF	143
3.3.3	Low-Dimensional Models	144
3.4	Inducing Word Embeddings from NLMs	150
3.5	Word Embedding Models	152
3.5.1	Word2Vec	152
3.5.2	GloVe	154
3.5.3	Remarks	158
3.6	Evaluating Word Embeddings	160
3.6.1	Extrinsic Evaluation	161
3.6.2	Intrinsic Evaluation	161
3.6.3	Visualization	165
3.7	Summary	165
4	Recurrent and Convolutional Sequence Models	169
4.1	Problem Statement	170
4.2	Recurrent Models	171
4.2.1	RNN-based Language Models	171
4.2.2	Training	173
4.2.3	Layer Stacking	176
4.2.4	Bi-directional Models	178
4.3	Memory	179
4.3.1	Memory as A System	180
4.3.2	Long Short-Term Memory	181
4.3.3	Gated Recurrent Units	183
4.4	Convolutional Models	185
4.4.1	Convolution	185
4.4.2	CNNs for Sequence Modeling	189
4.4.3	Handling Positional Information	191
4.5	Examples	196
4.5.1	Text Classification	196
4.5.2	End-to-End Speech Recognition	198
4.5.3	Sequence Labeling with LSTM and Graphical Models	201

4.5.4	Hybrid Models for Language Modeling	205
4.6	Summary	207
5	Sequence-to-Sequence Models	209
5.1	Sequence-to-Sequence Problems	210
5.2	The Encoder-Decoder Architecture	212
5.2.1	Encoding and Decoding	212
5.2.2	Example: Neural Machine Translation	214
5.3	The Attention Mechanism	217
5.3.1	A Basic Model	217
5.3.2	The QKV Attention	221
5.3.3	Multi-head Attention	224
5.3.4	Multi-layer Attention	227
5.3.5	Remarks	229
5.4	Search	234
5.4.1	The Length Problem	234
5.4.2	Pruning and Beam Search	238
5.4.3	Online Search	246
5.4.4	Exact Search	249
5.4.5	Differentiable Search	252
5.4.6	Hypothesis Diversity	254
5.4.7	Combining Multiple Models	256
5.4.8	More Search Objectives	258
5.5	Summary	261
6	Transformers	265
6.1	The Basic Model	265
6.1.1	The Transformer Architecture	265
6.1.2	Positional Encoding	269
6.1.3	Multi-head Self-attention	270
6.1.4	Layer Normalization	272
6.1.5	Feed-forward Neural Networks	273
6.1.6	Decoder-side Attention	274
6.1.7	Training and Inference	277
6.2	Syntax-aware Models	279
6.2.1	Syntax-aware Input and Output	280
6.2.2	Syntax-aware Attention Models	281
6.2.3	Multi-branch Models	283
6.2.4	Multi-scale Models	286
6.2.5	Transformers as Syntax Learners	287

6.3	Improved Architectures	291
6.3.1	Locally Attentive Models	291
6.3.2	Deep Models	295
6.3.3	Numerical Method-Inspired Models	301
6.3.4	Wide Models	304
6.4	Efficient Models	307
6.4.1	Sparse Attention	308
6.4.2	Recurrent and Memory Models	313
6.4.3	Low-dimensional Models	317
6.4.4	Parameter and Activation Sharing	322
6.4.5	Alternatives to Self-Attention	324
6.4.6	Conditional Computation	331
6.4.7	Model Transfer and Pruning	336
6.4.8	Sequence Compression	338
6.4.9	High Performance Computing Methods	340
6.5	Applications	343
6.5.1	Language Modeling	343
6.5.2	Text Encoding	344
6.5.3	Speech Translation	346
6.5.4	Vision Models	349
6.5.5	Multimodal Models	350
6.6	Summary	351

III

Large Language Models

7	Pre-training	361
7.1	Pre-training in NLP	362
7.1.1	Unsupervised, Supervised and Self-supervised Pre-training	362
7.1.2	Adapting Pre-trained Models	364
7.2	Self-supervised Pre-training Tasks	367
7.2.1	Decoder-only Pre-training	368
7.2.2	Encoder-only Pre-training	369
7.2.3	Encoder-Decoder Pre-training	376
7.2.4	Comparison of Pre-training Tasks	382
7.3	Example: BERT	384
7.3.1	The Standard Model	384
7.3.2	More Training and Larger Models	389
7.3.3	More Efficient Models	389
7.3.4	Multilingual Models	390

7.4	Applying BERT Models	392
7.5	Summary	398
8	Generative Models	399
8.1	A Brief Introduction to LLMs	400
8.1.1	Decoder-only Transformers	401
8.1.2	Training LLMs	404
8.1.3	Fine-tuning and Applying LLMs	405
8.1.4	Aligning LLMs with the World	411
8.1.5	Prompting LLMs	416
8.2	Training at Scale	421
8.2.1	Data Preparation	421
8.2.2	Model Modifications	423
8.2.3	Distributed Training	426
8.2.4	Scaling Laws	429
8.3	Long Sequence Modeling	432
8.3.1	Optimization from an HPC Perspective	433
8.3.2	Efficient Architectures	434
8.3.3	Cache and Memory	437
8.3.4	Sharing across Heads and Layers	446
8.3.5	Position Extrapolation and Interpolation	449
8.3.6	Remarks	460
8.4	Summary	462
9	Prompting	465
9.1	General Prompt Design	466
9.1.1	Basics	466
9.1.2	In-context Learning	469
9.1.3	Prompt Engineering Strategies	471
9.1.4	More Examples	476
9.2	Advanced Prompting Methods	486
9.2.1	Chain of Thought	486
9.2.2	Problem Decomposition	489
9.2.3	Self-refinement	496
9.2.4	Ensembling	502
9.2.5	RAG and Tool Use	506
9.3	Learning to Prompt	511
9.3.1	Prompt Optimization	512
9.3.2	Soft Prompts	515
9.3.3	Prompt Length Reduction	526
9.4	Summary	527

10	Alignment	529
10.1	Overview of LLM Alignment	530
10.2	Instruction Alignment	531
10.2.1	Supervised Fine-tuning	532
10.2.2	Fine-tuning Data Construction	537
10.2.3	Fine-tuning with Less Data	542
10.2.4	Instruction Generalization	543
10.2.5	Weak-to-Strong Generalization	545
10.3	Human Preference Alignment: RLHF	549
10.3.1	Basics of Reinforcement Learning	549
10.3.2	Training Reward Models	556
10.3.3	Training LLMs	559
10.4	Improved Human Preference Alignment	564
10.4.1	Better Reward Modeling	564
10.4.2	Direct Preference Optimization	571
10.4.3	Automatic Preference Data Generation	574
10.4.4	Step-by-step Alignment	576
10.4.5	Inference-time Alignment	579
10.5	Summary	580
11	Inference	583
11.1	Prefilling and Decoding	584
11.1.1	Preliminaries	584
11.1.2	A Two-phase Framework	589
11.1.3	Decoding Algorithms	592
11.1.4	Evaluation Metrics for LLM Inference	602
11.2	Efficient Inference Techniques	603
11.2.1	Prefix Caching	604
11.2.2	Batching	605
11.2.3	Parallelization	614
11.2.4	Remarks	615
11.3	Inference-time Scaling	616
11.3.1	Context Scaling	617
11.3.2	Search Scaling	618
11.3.3	Output Ensembling	619
11.3.4	Thinking-Path Generation and Verification	620
11.4	Summary	628
12	Reasoning	631
12.1	Background and Concepts	631
12.1.1	What is Reasoning?	631

12.1.2	Why Reasoning Matters	633
12.1.3	From LLMs to Reasoning Models	634
12.2	Reasoning via Prompting	635
12.3	Inference-time Reasoning	636
12.3.1	Inference-Time Compute Scaling	636
12.3.2	Sampling Multiple Reasoning Paths	638
12.3.3	Search Strategies for Reasoning	641
12.3.4	Verification	642
12.4	Learning to Reason	646
12.4.1	Supervised Reasoning Data	647
12.4.2	Process Supervision	648
12.4.3	On-Policy Distillation for Reasoning	652
12.4.4	Scaling Training for Reasoning	655
12.5	Evaluating Reasoning LLMs	656
12.5.1	Outcome-based vs. Process-based Evaluation	656
12.5.2	Evaluation Metrics	657
12.5.3	Open Issues	659
12.6	Remarks	659
12.6.1	External Reasoning vs. Internal Reasoning	660
12.6.2	Memorization vs. Reasoning	661
12.6.3	Can LLMs Truly Reason?	663
12.7	Summary	664